



Data Science on AWS ×

▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker

Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines

Contributors

Summary

[Data Science on AWS](#) > [Introduction](#) > [Prerequisites](#)

Prerequisites

What do I need to run this workshop?

- You need a computer or laptop with a browser, preferably with the latest version of Chrome / FireFox.
- Sequentially read and follow the instructions described in [AWS Hosted Event](#) and [Workshop Setup](#).

What background knowledge is necessary for this workshop?

It will be easier for you to run this workshop if you have:

- Experience with deep learning models
- Familiarity with Python or other similar programming languages
- Experience with Jupyter notebooks
- Beginners level knowledge and experience with Amazon SageMaker

Who is this workshop for?

- This workshop is designed for data and ML practitioners who want to learn how to

Data Science on AWS



▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker

Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines

Contributors


Summary

processing (NLP) and generative AI in the AWS cloud.

How long does it take to complete this workshop?

- Approximately 4 hours.

Will this workshop cost me anything?

Going through this workshop outside an AWS hosted event will incur some cost. Make sure you clean up your resources at the end to avoid unnecessary charges. See [AWS Pricing](#)  for more details.

Previous

Next