

**Data Science on AWS** ✕

▼ Introduction

[Prerequisites](#)[AWS Hosted Event](#)[Workshop Setup](#)[Part 1: Distributed Data](#)[Preprocessing](#)[Part 2: Fine-tuning LLMs with Amazon SageMaker](#)[Part 3: Automating Fine-Tuning workflows with SageMaker](#)[Pipelines](#)[Contributors](#)[Summary](#)[Data Science on AWS](#) > **Part 2: Fine-tuning LLMs with Amazon SageMaker**

Part 2: Fine-tuning LLMs with Amazon SageMaker

Learn how to build generative AI and large language models (LLMs) for your data science projects that support both interactive development and ephemeral training jobs.

In this workshop, we will walk through a series of hands-on labs.

Instructions

In SageMaker Studio, open and run the following notebooks in the sequence shown here:

0. Lab overview
1. Setup workshop dependencies
5. Analyze the impact of prompt engineering using a HuggingFace model
6. Perform feature engineering on a raw text dataset using HuggingFace
 - Option A: Notebook processing in SageMaker Studio
 - Option B: SageMaker Processing Job
7. Fine-tune a HuggingFace model for dialogue summarization

Data Science on AWS



▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

**Part 2: Fine-tuning LLMs with
Amazon SageMaker**

Part 3: Automating Fine-Tuning
workflows with SageMaker

Pipelines

Contributors

Summary

◦ Option B: SageMaker Training Job



Go to the provided AWS account.

Run the notebooks in SageMaker Studio.

Previous

Next