

**Data Science on AWS** ✕

▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker**Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines**

Contributors

Summary

[Data Science on AWS](#) > **Part 3: Automating Fine-Tuning workflows with SageMaker Pipelines**

Part 3: Automating Fine-Tuning workflows with SageMaker Pipelines

Learn how to build repeatable MLOps pipelines for your data science projects including prompt engineering, generative AI, and large-language models (LLMs).

In this workshop, we will walk through a series of hands-on labs.

Instructions

In SageMaker Studio, open and run the following notebooks in the sequence shown here:

0. Lab overview
1. Setup workshop dependencies
-
-
-
-
-
8. Create an automated end-to-end ML MLOps workflow with SageMaker Pipelines
9. Deploy a fine-tuned generative AI model to a real-time SageMaker Endpoint
10. Run inference on a SageMaker Endpoint in real time

**Go to the provided AWS account.**

Run the notebooks in SageMaker Studio.

[Previous](#)[Next](#)

Data Science on AWS ×

▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker

**Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines**

Contributors

Summary