



Data Science on AWS ×

▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data

Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker

Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines

Contributors

Summary

Data Science on AWS > Contributors

Contributors

Antje Barth



Antje is a **Principal Developer Advocate for AI and machine learning** at AWS. She is co-author of the O'Reilly book "Data Science on AWS." Antje frequently speaks at AI/ML conferences, events, and meetups around the world. She also co-founded the Düsseldorf chapter of Women in Big Data.

Chris Fregly



Chris is a **Principal Specialist Solution Architect for AI and**

Data Science on AWS



▼ Introduction

Prerequisites

AWS Hosted Event

Workshop Setup

Part 1: Distributed Data
Preprocessing

Part 2: Fine-tuning LLMs with
Amazon SageMaker


Part 3: Automating Fine-Tuning
workflows with SageMaker
Pipelines

Contributors

Summary

AWS." Chris is also the Founder of many global meetups focused on Apache Spark, TensorFlow, Ray, and KubeFlow. He regularly speaks at AI and machine learning conferences across the world including O'Reilly AI, Open Data Science Conference, and Big Data Spain.

Clay Elmore



Clay is a **Specialist Solution Architect for AI and machine learning** at AWS. After spending many hours in a materials research lab, his background in chemical engineering was quickly left behind to pursue his interest in machine learning. He has worked on ML applications in many different industries ranging from energy trading to hospitality marketing. Clay has a special interest in bringing software development practices to ML and guiding customers towards repeatable, scalable solutions by using these principles.

Previous

Next